

Victor Giannakouris

victorasgs@gmail.com • (+30)698 4384424 • <http://github.io/gsvic>

EDUCATION

Ionian University, Department of Informatics, Corfu, Greece

- Diploma in Informatics (4 years degree; 240 ECTS) Sep 2012- May 2017
 - Thesis: : Optimization of Distributed SQL Query Execution Over Multiple-Engine Environments
 - Advisor: Prof. Dimitrios Tsoumakos

EXPERIENCE

Unravel Data Systems, Palo Alto, CA, US

- Software Engineer Oct 2018 – Present
 - Big Data Application Performance Management (APM)

Stealth Mode Silicon Valley Startup, Palo Alto, CA, US

- Data Engineer Apr 2018 – Sep 2018
 - Development of data collection microservices on the company's platform's data ingestion component.
 - Tools: Kubernetes, AWS, Google Cloud Platform, Java, Python

BEAT, Athens

- Big Data Software Engineer Sep 2017 – Apr 2018
 - Development of streaming and ETL jobs over large-scale data.
 - Tuning and optimization of Spark SQL batch jobs on query and storage level.
 - Tools: Spark, YARN, Hadoop, Hive, Presto, Scala, Python

Computing Systems Laboratory (CSLab), NTUA, Greece

- Researcher Jul 2015 – Aug 2017
 - R&D on Big Data Management, Databases and Multi-Engine Analytics.
 - Participated in four publications.
 - Development of a middleware optimization framework for Spark SQL, achieving execution time improvement for queries over heterogeneous environments up to one order of magnitude.
 - Participated in the EU Funded Project ASAP(A Scalable Analytics Platform).

isMOOD Data Technology Services

- Software Engineer Apr 2014 – Sep 2015
 - Data collection and integration into a unified source
 - Development, scheduling & optimization of web crawlers/spiders
 - Tools: Scrapy, Python, Elasticsearch, MongoDB

AWARDS

- IEEE BigData 2016, Student Travel Award 2016
 - Won one of the 36 given.
 - Published two full-papers (10 pages) in the research-track.

PRESENTATIONS - TALKS

- Hellenic Data Management Symposium (HDMS) 2017, Athens, Greece Aug 2017
Distributed SQL Query Execution over Mutli-Engine Environments.
- IEEE Big Data 2016, Washington, DC, USA Dec 2016
Distributed SQL Query Execution over Mutli-Engine Environments.

SKILLS

- **Programming Languages:** Java, Scala, Python, SQL
- **Databases:** Hive, Presto, Elasticsearch, MongoDB, MemSQL, MariaDB, PosgreSQL
- **Data Management - Analytics:** Spark, Hadoop, scikit-learn
- **Infrastructure:** Kubernetes, AWS, Openstack
- **Software Development:** Agile, Scrum

CERTIFICATIONS

- SQL (Mini-Course) Stanford Online
- CS190.1x: Scalable Machine Learning edX - BerkeleyX
- CS100.1x: Introduction to Big Data with Apache Spark edX - BerkeleyX
- MAS.S69x: Big Data and Social Physics edX - MITx
- HAISS 2013 – LEARNING (AI Summer School) EETN - Ionian University

PUBLICATIONS JOURNALS

- [1] X. Evangelopoulos, V. Giannakouris, L. Iliadis, C. Makris, Y. Plegas, A. Plerou, S. Sioutas, "Evaluating Information Retrieval using Document Popularity: A case study using MapReduce," *Engineering Applications of Artificial Intelligence*, vol. 89, no. 2, pp. 4123–4133, Feb 2008.

CONFERENCES

- [1] V. Giannakouris, N. Papailiou, D. Tsoumakos, and N. Koziris, “MuSQLE: Distributed SQL Query Execution Over Multiple-Engine Environments,” *In proceedings of the 2016 IEEE International Conference on Big Data (BigData 2016)*, Washington, DC, USA, Dec 2016.
- [2] K. Doka, N. Papailiou, V. Giannakouris, D. Tsoumakos and N. Koziris, “Mix ’n Match Multi-Engine Analytics,,” *In proceedings of the 2016 IEEE International Conference on Big Data (BigData 2016)*, Washington, DC, USA, Dec 2016.
- [3] K. Doka, M. Filatov, V. Giannakouris, V. Kantere, N. Koziris, C. Mantas, N. Papailiou, V. Papaioannou and D. Tsoumakos, “Optimizing, Planning and Executing Analytics Workflows over Multiple Engines.,” *In the 1st International Workshop on Multi-Engine Data Analytics (MEDAL 2016)*, Bordeaux, France, Dec 2016.
- [4] N. Papailiou, K. Doka, V. Giannakouris, V. Papaioannou, D. Tsoumakos and N. Koziris, “Robust and Adaptive Multi-Engine Analytics using IReS (demo paper),” *In the 10th International Workshop on Enabling Real-Time Business Intelligence (BIRTE 2016), in conjunction with the VLDB Conference*, New Delhi, India, Mar 2016.
- [5] V. Giannakouris, A. Plerou, and S. Sioutas, “CSMR: A Scalable Algorithm for Text Clustering with Cosine Similarity and MapReduce,” *In the 10th International Conference on Artificial Intelligence Applications and Innovations*, Rhodes, Greece, Sep 2014.

LANGUAGES

- **English** (fluent), **Greek** (native)